

Detecting and Mitigating Semantic Drift in Multi-Turn Instruction-Based Image Editing

R. W. V. Imansha Dilshan
Software Engineering Teaching Unit
University of Kelaniya
rajapak-se21029@stu.kln.ac.lk

Abstract—Instruction-based image editors such as InstructPix2Pix let users apply a sequence of natural-language edits to an image, yet no standard metric captures what happens to the parts of the image nobody asked to change. This paper introduces the Drift Score, a model-agnostic metric that combines Segment Anything (SAM) region segmentation with CLIP embedding similarity to quantify unintended change after an edit. The metric is applied across 120 four-step edit chains built from 60 COCO images, split evenly between localized object-level edits and broad scene-level edits. Two low-cost mitigation strategies, masked conditioning and region-locking, are implemented and evaluated against an unmitigated baseline using paired significance testing. Both mitigations produce large, highly significant reductions in measured drift (region-locking: 89.1%, $p < 0.0001$; masked conditioning: 68.6%, $p < 0.0001$). Two pre-registered hypotheses did not hold as predicted: region-locking significantly outperforms masked conditioning rather than the reverse, and drift does not increase across the chain, with the only significant step-position effect running opposite to the expected compounding pattern. Both reversals are traced to specific mechanisms rather than left unexplained. The Drift Score, dataset, and full pipeline are released as a reproducible toolkit built entirely from pretrained models.

Index Terms—Image Editing, Diffusion Models, Semantic Drift, CLIP, Segment Anything, Evaluation Metrics, Instruction-Based Editing

I. INTRODUCTION

Iterative refinement is a natural way to use an instruction-based image editor. A user might ask the model to make the sky look like sunset, then add a bird, then remove a fence in the background, refining an image through several small edits rather than one large prompt. This pattern introduces a failure mode that single-step evaluation cannot see: a model can follow every individual instruction correctly while still drifting further and further from what the user actually wanted, simply because nothing in its training or inference constrains it to leave the rest of the image alone. Existing evaluation protocols for these models, including CLIP-based instruction-following scores and human preference studies, measure whether the requested change happened. None of them measure what else changed as a side effect. This paper calls that side effect *semantic drift*.

Most existing instruction-based editors, InstructPix2Pix chief among them [1], are trained on synthetically generated (image, instruction, edited-image) triplets. That training recipe supplies no explicit signal for “leave everything else untouched” beyond whatever the data-generation process happened to preserve. Drift is therefore an expected consequence

of how these models are built, not merely an empirical curiosity worth measuring after the fact.

Three questions follow directly from this observation. Can unintended drift be measured automatically, without manual inspection of every edited image (RQ1)? Does drift accumulate as a chain of edits gets longer, so that later instructions cause more collateral damage than earlier ones (RQ2)? And can cheap, pretrained-model-only interventions reduce it (RQ3)? This paper answers all three.

The contribution is threefold. First, the Drift Score, a reusable metric combining Segment Anything (SAM) region segmentation with CLIP embedding similarity to quantify unintended regional change, at both the single-edit and full-chain level. Second, an empirical study of how drift behaves across 120 real four-step edit chains, split between localized object-level edits and broad scene-level edits. Third, a statistically grounded comparison of two mitigation strategies, masked conditioning and region-locking, against an unmitigated baseline.

The work was scoped for a single contributor on a 30-day timeline using only pretrained models and free-tier compute: a laptop for development and analysis, and a free Google Colab GPU for the heavier inference stages. No model was trained or fine-tuned. The contribution here is the measurement methodology and the empirical findings it produces, not a new editing model.

II. RELATED WORK

A. Foundations

InstructPix2Pix [1] is the editor evaluated throughout this study. It fine-tunes a latent diffusion model [2] on a large corpus of synthetically generated editing triplets, produced by combining a language model with Prompt-to-Prompt-style attention injection [3]. Because supervision comes entirely from this generated data, the model inherits no explicit signal for preserving untouched content beyond what that generation process happened to keep stable, which is the mechanistic root of the drift problem this paper measures. CLIP [4] supplies the embedding space the Drift Score measures change in, and the text-image similarity used to identify which region an instruction targets. Segment Anything [5] supplies the region decomposition without which drift could only be measured at whole-image granularity, obscuring exactly the localized side effects this paper is concerned with.

B. Consistency-Preserving Editing

Prompt-to-Prompt [3], Plug-and-Play [6], and MasaCtrl [7] are generation-time techniques for keeping an edited image consistent with its source, by manipulating cross-attention maps, injecting spatial features, or converting self-attention into a mutual query against the source image’s own generation trajectory. All three aim to *produce* a more consistent single edit. None of them measure how much unintended change remains afterward, and none address a chain of sequential edits. MasaCtrl’s own motivating discussion states the semantic-drift problem in almost the same terms used here, but stops at fixing one edit rather than measuring or chaining it.

C. Multi-Turn Editing Data

MagicBrush [8] already contains manually annotated multi-turn editing sessions, which is independent validation that a four- or five-step chain is a realistic usage pattern rather than an artificial stress test. Its evaluation protocol, however, only checks whether each turn’s target region changed the way it was supposed to. It never measures how much everything else changed in the process, which is exactly the gap this paper addresses.

TABLE I
CAPABILITY COMPARISON OF RELATED APPROACHES

Approach	Unintended Change	Chain Aware	Mitigation Compare	Model Agnostic
Prompt-to-Prompt [3]	×	×	×	×
Plug-and-Play [6]	×	×	×	×
MasaCtrl [7]	×	×	×	×
MagicBrush [8]	Partial	✓	×	×
Proposed (Drift Score)	✓	✓	✓	✓

No prior approach measures unintended regional change across a chain of edits or compares mitigation strategies on that specific axis (Table I); these dimensions are author-defined for this comparison rather than drawn from a standardized benchmark. The contribution of this paper is that combination: a measurement instrument built on top of the segmentation and embedding tools the field already uses, applied to the chain-based usage pattern the field already knows is realistic.

III. METHODOLOGY

The pipeline realizes a Segment → Identify → Edit → Score cycle at every step of a chain (Fig. 1), repeated once per instruction.

A. The Drift Score

Given a pre-edit image, a post-edit image, and the instruction that produced it, the pre-edit image is first segmented into regions with SAM, and each region is cropped from both the pre- and post-edit image at the *same* coordinates. Segmenting the post-edit image independently would yield unrelated region identifiers with no correspondence between the two images, so the same box set is deliberately reused. The target region, the one box the instruction most plausibly refers to, is identified by embedding every region crop and the instruction text with




Fig. 1. Per-step pipeline. The mitigation branch is absent for the unmitigated baseline and active for the two mitigation strategies.

CLIP and taking the highest cosine similarity. For every non-target region r , the per-region drift is

$$d_r = 1 - \cos(\text{CLIP}(x_r^{\text{pre}}), \text{CLIP}(x_r^{\text{post}})) \quad (1)$$

and the per-step Drift Score is the mean of d_r over all non-target regions. A chain’s cumulative Drift Score is the sum of its per-step scores:

$$D_{\text{chain}} = \sum_{t=1}^T \frac{1}{|R_t|} \sum_{r \in R_t} d_{r,t} \quad (2)$$

where R_t is the set of non-target regions at step t and T is the chain length. Higher values indicate more unintended change accumulated across the chain.

B. Dataset

Sixty images were sampled from COCO val2017 (seed 42, reproducible), filtered to those with two to eight annotated objects, enough to guarantee both a plausible target region and several non-target regions without an overly cluttered scene. Every image was inspected individually and given two hand-written, per-image instruction chains rather than instructions drawn from a fixed template pool:

- **Chain A (object-level):** four localized, single-object edits, for example changing a horse’s color and removing an obstacle in front of it.
- **Chain B (global):** four broad, scene-level edits, for example making a scene rainy and adding sunset lighting.

This produced 120 chains in total: 60 object-level and 60 global, each four instructions long.

C. Pipeline and Models

TABLE II
MODELS USED, ALL PRETRAINED AND PUBLICLY AVAILABLE

Component	Model	Role
Editor	InstructPix2Pix	Executes each edit, 512×512
Segmentation	SAM (ViT-B)	Region decomposition
Embedding	CLIP (ViT-B/32)	Change measurement, targeting

No model was trained or fine-tuned. Heavy inference, editing and segmentation, ran on a free-tier Colab GPU; drift-score aggregation and statistical analysis ran locally on CPU.

D. Mitigation Strategies

Two strategies were implemented, both identifying the target region through the identical SAM-plus-CLIP procedure the Drift Score itself uses, so that measurement and mitigation agree on what counts as the target.

Region-locking runs the edit on the full image as usual, then reverts every pixel outside the target region’s box back to the pre-edit image. This is a correction applied after generation.

Masked conditioning crops the image down to just the target region, with a 32-pixel padding margin included for surrounding context, before generation. The edit is applied only to that crop, and the result is pasted back into the full frame. The model never sees the rest of the image, which makes this a constraint applied at generation time rather than afterward.

A third strategy considered during scoping, attention-restricted editing informed by Prompt-to-Prompt-style cross-attention control, was treated as optional and was not implemented, since RQ3 already has a clear, statistically supported answer without it.

E. Statistical Analysis

Paired t -tests and Wilcoxon signed-rank tests compare cumulative Drift Scores between conditions, paired by image identity and chain type. The Wilcoxon test is treated as the primary result given the small, likely skewed sample; the t -test is reported alongside it for completeness. Of 480 baseline edit steps, 13 (2.7%) could not be scored, because SAM occasionally finds too few regions in a heavily stylized pre-edit image for anything to remain once the target region is excluded. These were dropped from pairwise comparisons rather than imputed.

IV. EXPERIMENT AND EVALUATION

A. Experimental Setup

All 120 chains were run three times: once unmitigated (baseline), once with region-locking active at every step, and once with masked conditioning active at every step. Each run produced five images per chain, the original plus one output per instruction, all committed alongside their Drift Scores for reproducibility. The baseline chain outputs also served as the pre-edit input for the corresponding mitigated chains at each step, so all three conditions edit from the same starting point at every stage.

B. RQ1: Validating the Drift Score

Before trusting the score at scale, chains spanning the full score range were inspected by eye. The lowest-scoring baseline chain (0.082, a surfer riding a wave) shows only the requested outfit and surfboard color changes; the wave, spray, and sky are the same scene, pixel for pixel, outside those two objects. The highest-scoring baseline chain (0.565, a boy holding a baseball glove) is a genuine catastrophic failure: by the third step the image has collapsed to near-total black after the second edit destroyed the original composition, and the fourth instruction is applied on top of the wreckage (Fig. 3,

second panel). The score correctly identifies this as the single worst chain in the dataset, which is the kind of agreement between the metric and visible content that a purely numerical validation cannot substitute for.

C. RQ3: Mitigation Effectiveness

Both mitigations reduce mean cumulative drift by a wide margin relative to the unmitigated baseline (Table III, Fig. 2), and the reduction is significant well beyond conventional thresholds in every breakdown tested.

TABLE III
CUMULATIVE DRIFT SCORE BY CONDITION ($n = 120$)

Condition	Mean	Reduction	Wilcoxon p
Baseline	0.293	—	—
Masked Conditioning	0.092	68.6%	< 0.0001
Region-Locking	0.032	89.1%	< 0.0001




Fig. 2. Mean cumulative Drift Score by condition. Error bars omitted; both reductions are significant at $p < 0.0001$ by paired t -test and Wilcoxon signed-rank test.

Region-locking also outperforms masked conditioning directly, in a head-to-head paired comparison ($p < 0.0001$). Broken down by chain type (Table IV), both mitigations perform slightly better on global chains than on object-level chains, though both remain highly significant in each subgroup, suggesting the target-identification weak point discussed in Section V-C, which is specific to object-adding instructions, affects object-level chains more than global ones. On the worst-case chain from Section IV-B, baseline cumulative drift of 0.565 falls to 0.076 under region-locking and 0.180 under masked conditioning, both large, visually confirmed recoveries (Fig. 3).

TABLE IV
DRIFT REDUCTION BY CHAIN TYPE ($n = 60$ PER CELL)

Chain Type	Condition	Reduction	Wilcoxon p
Object-Level	Masked Conditioning	62.5%	< 0.0001
Object-Level	Region-Locking	86.0%	< 0.0001
Global	Masked Conditioning	74.8%	< 0.0001
Global	Region-Locking	92.2%	< 0.0001




Fig. 3. The worst-case baseline chain (baseball glove) under all three conditions. The baseline collapses to near-black; region-locking recovers the scene almost exactly, with one visible seam artifact near the glove; masked conditioning recovers the scene for three of four steps but produces an unrelated image on the final “add a cap” instruction (Section V-A).

D. RQ2: Compounding Across the Chain

No single consecutive step transition is statistically significant (step 1 to 2: $p = 0.85$; 2 to 3: $p = 0.33$; 3 to 4: $p = 0.47$), meaning there is no detectable escalation at any individual point in the chain (Fig. 4). Comparing the two endpoints directly, step 1 versus step 4 is significant ($p = 0.012$ by t -test, $p = 0.006$ by Wilcoxon), but with step 4 showing *less* drift (mean 0.066) than step 1 (mean 0.080), the reverse of the predicted direction.




Fig. 4. Mean baseline Drift Score by chain position. The step 1 to step 4 difference is significant ($p = 0.012$); no individual step-to-step transition is.

V. DISCUSSION

A. Why the Compounding Hypothesis Did Not Hold

The working hypothesis going into this study was that later edits would show progressively more collateral damage, because each one operates on an already-modified image. The data shows the opposite, at the one point where a significant difference exists at all. The most plausible explanation, consistent with the baseball-glove example above, is that a chain which collapses catastrophically early has, by construction of a CLIP-similarity metric, very little room left to register as further changed in subsequent steps. An image that is already almost entirely black cannot become much more different from its original self than it already is. This is a measurement ceiling effect masking a front-loaded failure, not evidence that later edits are genuinely gentler. Separating “drift is front-loaded and severe” from “drift does not compound” would need a metric sensitive to further degradation of an already-degraded

image, which is a natural direction for follow-up work rather than something the CLIP-similarity approach used here can resolve as built.

B. Why Region-Locking Outperformed Masked Conditioning

The prediction going in was that masked conditioning, a stricter constraint applied at generation time, would outperform region-locking’s after-the-fact correction, at some cost to edit quality at the region boundary. The opposite happened on the drift metric itself, and the likely reason is a property of how the metric and the mitigation interact rather than either mitigation being fundamentally weaker. Masked conditioning edits a padded crop, the target box plus a 32-pixel margin, deliberately included so the model has enough surrounding context to produce a coherent edit. The Drift Score, however, only excludes the raw target box from scoring, so masked conditioning’s own padding ring is counted as drift even when nothing objectionable happened there. Region-locking’s hard revert has no equivalent margin. This should not be read as proof that region-locking produces better-looking edits in general; the qualitative example in Fig. 3 shows region-locking has its own artifact, a visible, unblended rectangular seam near the glove, consistent with the boundary-cost intuition behind the original hypothesis. It simply does not show up as a drift-score cost the way that hypothesis predicted.

C. A Shared Weak Point: Target-Region Identification

Both mitigations, and the Drift Score itself, depend on CLIP correctly identifying which region an instruction targets. This is reliable for “change X” and “remove X” instructions, where the target already exists in the pre-edit image. It is not reliable for “add X” instructions, where nothing pre-existing can match well. Tracing one concrete failure directly: for the instruction “add a cap on the boy’s head,” the final step of the worst-case chain, the identified target box covered 99.6% of the frame. Masked conditioning’s spatial constraint therefore gave no real protection for that step, since it was cropping to almost the entire image regardless. Region-locking faced the identical oversized box on the same step and produced a fine result anyway, most likely because diffusion sampling is stochastic with no fixed seed, so the two mitigations’ generation calls for that step were independent random draws from the same model, and one landed conservative while the other did not. This is a reproducible limitation of the target-identification approach, not an implementation bug in either mitigation, and it affects roughly half of the object-level chains, those whose instructions use an “add” verb.

To check whether this single traced case reflects a dataset-wide pattern, every step across all 480 was split into “add” instructions and all others, and compared with an independent-samples Mann-Whitney U test, since add and non-add steps are not paired observations within the same chain (Table V). Baseline and region-locking both show lower mean drift on add steps than on other steps, a difference that is only borderline significant ($p = 0.052$ and $p = 0.072$ respectively) rather than conclusive. Masked conditioning is the outlier: its

add-step mean is numerically *higher* than its other-step mean, the only condition of the three where this reverses, though the difference itself does not reach significance ($p = 0.693$) at this sample size. The direction is consistent with the mechanism traced above, an oversized target box removing masked conditioning’s spatial constraint specifically on add instructions, but the dataset-wide test should be read as corroborating, not confirming, the single traced example; a larger sample would be needed to establish the effect conclusively.

TABLE V
PER-STEP DRIFT BY INSTRUCTION TYPE

Condition	“Add” Mean	Other Mean	Mann-Whitney p
Baseline	0.068	0.078	0.052
Region-Locking	0.006	0.009	0.072
Masked Conditioning	0.026	0.022	0.693

D. Limitations

Six constraints bound the present study. The 60-image, 120-chain dataset is adequate for the paired tests reported here but should be read as an initial empirical study rather than a definitive benchmark. All results are specific to InstructPix2Pix; the Drift Score is designed to be model-agnostic, but that design goal was not tested empirically against a second editor within this project’s scope. CLIP similarity is an imperfect proxy for human-perceived change; an interactive human-rating tool covering 15 stratified chains was built for this purpose but was not administered, an explicitly optional step skipped under the 30-day time budget, so the CLIP-based numbers here carry that caveat without an independent human check. Target-region identification through CLIP’s top-1 similarity is a known weak point for “add” instructions, as shown above. Region boxes, not exact per-pixel segmentation masks, are used throughout for both scoring and mitigation, for simplicity and consistency, and this coarser spatial unit is the direct cause of masked conditioning’s padding-margin effect. Finally, the attention-restricted editing mitigation considered during scoping was not implemented, since RQ3 has a clear answer without it.

E. Reproducibility

The full pipeline is organized as a small, dependency-light toolkit rather than a single notebook. Drift-score computation, region segmentation, and CLIP embedding each live in their own module with unit tests that exercise the scoring logic against synthetic vectors and mocked segmentation output, independent of any GPU or network access. The two mitigation strategies share the same target-identification code path as the scorer by construction, rather than by convention, which is what makes the mechanistic explanations in Sections V-B and V-C traceable to a specific function rather than to inferred behavior. Every image, edit chain, intermediate result, and statistical test output is version-controlled alongside the code that produced it, so any reported number in this paper can be regenerated from the committed dataset without re-running the GPU stages. The chain-construction and scoring approach is

not specific to InstructPix2Pix; substituting a different editor requires only replacing the module that calls it.

F. Future Work

The most direct extension is evaluating the Drift Score against a second, more recent editing model, to test the model-agnosticism claimed but not empirically exercised here. Running the built human-perceptual check would validate the CLIP-based drift numbers against human judgment on the same 15 chains. A drift metric less prone to the ceiling effect identified in Section V-A, one that can register further degradation of an already-degraded image, would let a future study separate gradual compounding from sudden, front-loaded failure with more confidence than the present metric allows.

VI. CONCLUSION

This paper set out to answer whether unintended drift across a chain of instruction-based image edits can be measured automatically, whether it compounds, and whether cheap mitigations reduce it. RQ1 and RQ3 have clear, positive, and statistically well-supported answers: yes, and yes, by a wide margin, with region-locking cutting mean cumulative drift by 89.1% and masked conditioning by 68.6%, both at $p < 0.0001$. RQ2’s answer is more interesting than a simple yes or no. The data suggests drift may be front-loaded into occasional catastrophic failures rather than gradually compounding, a pattern the Drift Score as built cannot fully distinguish from no compounding at all. Two specific mechanisms, a measurement ceiling effect and a metric-mitigation margin mismatch, explain why two of the original hypotheses reversed rather than simply failing to replicate, which is itself evidence that the underlying measurement is capturing something real and interpretable rather than noise.

Two constraints bound these conclusions. All results are specific to InstructPix2Pix, and the human-perceptual validation step was built but not administered. Resolving both, together with testing the metric against a second editing model, constitutes the next step toward a benchmark the field can rely on rather than a single-project case study.

ACKNOWLEDGMENT

This work was completed independently within a 30-day, free-tier compute budget. The author thanks Google Colab for free GPU access, and the respective research groups behind InstructPix2Pix, Segment Anything, and CLIP for releasing pretrained checkpoints publicly, without which a project of this scope would not have been feasible for a single undergraduate contributor.

REFERENCES

- [1] T. Brooks, A. Holynski, and A. A. Efros, “InstructPix2Pix: Learning to Follow Image Editing Instructions,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [2] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, “High-Resolution Image Synthesis with Latent Diffusion Models,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2022.

- [3] A. Hertz, R. Mokady, J. Tenenbaum, K. Aberman, Y. Pritch, and D. Cohen-Or, "Prompt-to-Prompt Image Editing with Cross Attention Control," *arXiv preprint arXiv:2208.01626*, 2022.
- [4] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning Transferable Visual Models From Natural Language Supervision," in *Proc. Int. Conf. Machine Learning (ICML)*, 2021.
- [5] A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, P. Dollár, and R. Girshick, "Segment Anything," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2023.
- [6] N. Tumanyan, M. Geyer, S. Bagon, and T. Dekel, "Plug-and-Play Diffusion Features for Text-Driven Image-to-Image Translation," in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [7] M. Cao, X. Wang, Z. Qi, Y. Shan, X. Qie, and Y. Zheng, "MasaCtrl: Tuning-Free Mutual Self-Attention Control for Consistent Image Synthesis and Editing," in *Proc. IEEE/CVF Int. Conf. Computer Vision (ICCV)*, 2023.
- [8] K. Zhang, L. Mo, W. Chen, H. Sun, and Y. Su, "MagicBrush: A Manually Annotated Dataset for Instruction-Guided Image Editing," in *Adv. Neural Inf. Process. Syst. (NeurIPS), Datasets and Benchmarks Track*, 2023.